

TensorFlow: a Framework for Scalable Machine Learning

ACM Learning Center, 2016

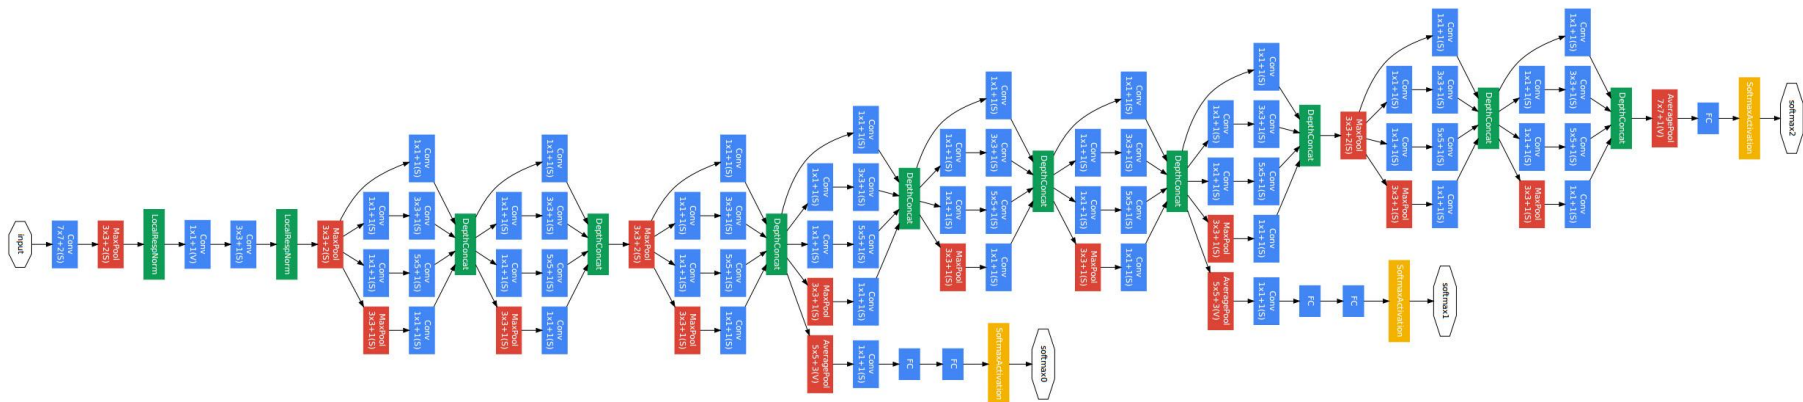
You probably want to know...

- What is TensorFlow?
- Why did we create TensorFlow?
- How does TensorFlow work?
- Code: Linear Regression
- Code: Convolution Deep Neural Network
- Advanced Topics: Queues and Devices



- Fast, flexible, and scalable open-source machine learning library
- One system for research and production
- Runs on CPU, GPU, TPU, and Mobile
- Apache 2.0 license


Modeling complexity



Machine learning gets complex quickly




**Distributed
System**



**Heterogenous
System**

TensorFlow Handles Complexity



Modeling complexity



**Distributed
System**



**Heterogenous
System**

Under the Hood

The background is a solid orange color. Overlaid on this is a white wireframe graphic of a mountain range. The mountains are composed of a grid of lines that form a series of peaks and valleys, creating a three-dimensional effect. The text "Under the Hood" is centered in the middle of the image in a white, bold, sans-serif font.

A multidimensional array.



TensorFlow




A graph of operations.

The TensorFlow Graph

Computation is defined as a graph

- Graph is defined in high-level language (Python)
- Graph is compiled and optimized
- Graph is executed (in parts or fully) on available low level devices (CPU, GPU, TPU)
- Nodes represent computations and state
- Data (tensors) flow along edges



Build a graph; then run it.


...

```
c = tf.add(a, b)
```


...

```
session = tf.Session()
```

```
value_of_c = session.run(c, {a=1, b=2})
```

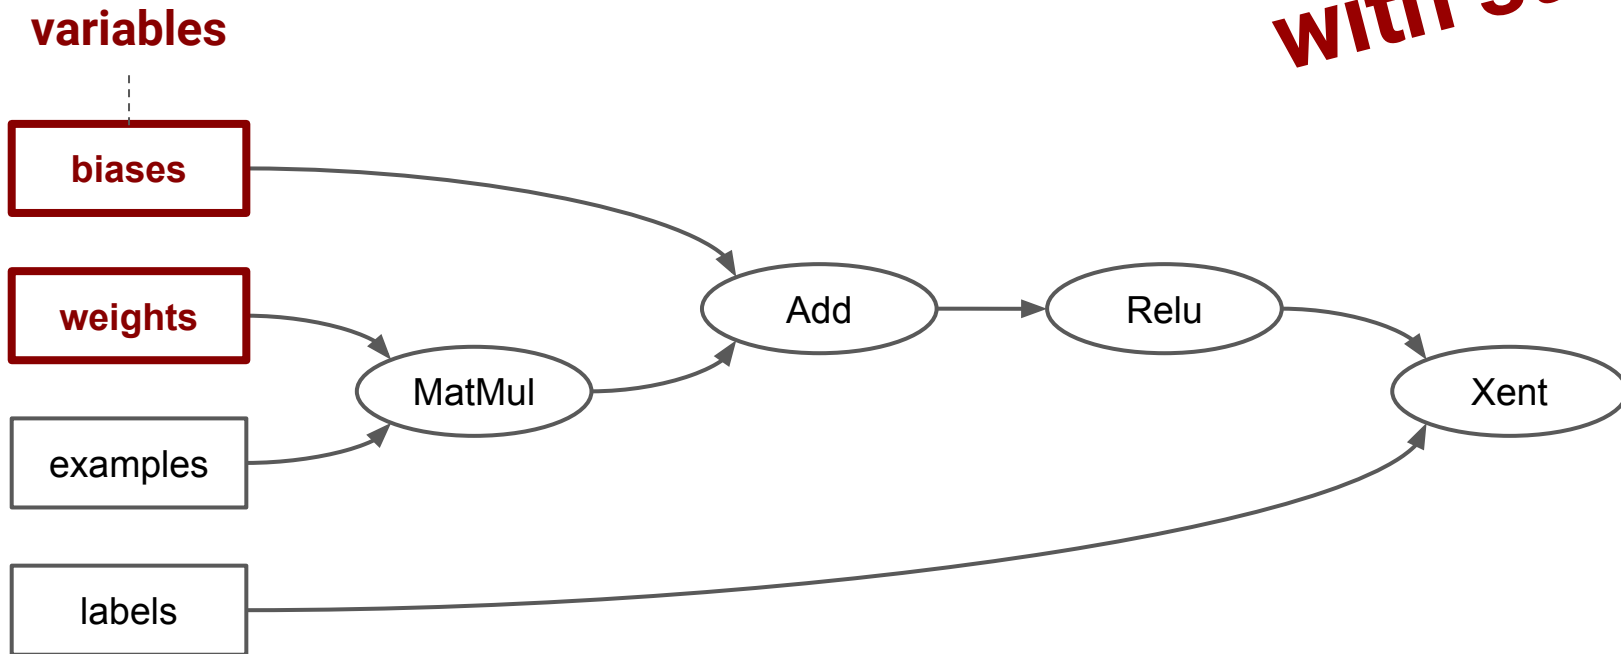


Any Computation is a TensorFlow Graph




Any Computation is a TensorFlow Graph

with state



Automatic Differentiation


**Automatically add ops which
compute gradients for variables**



Any Computation is a TensorFlow Graph

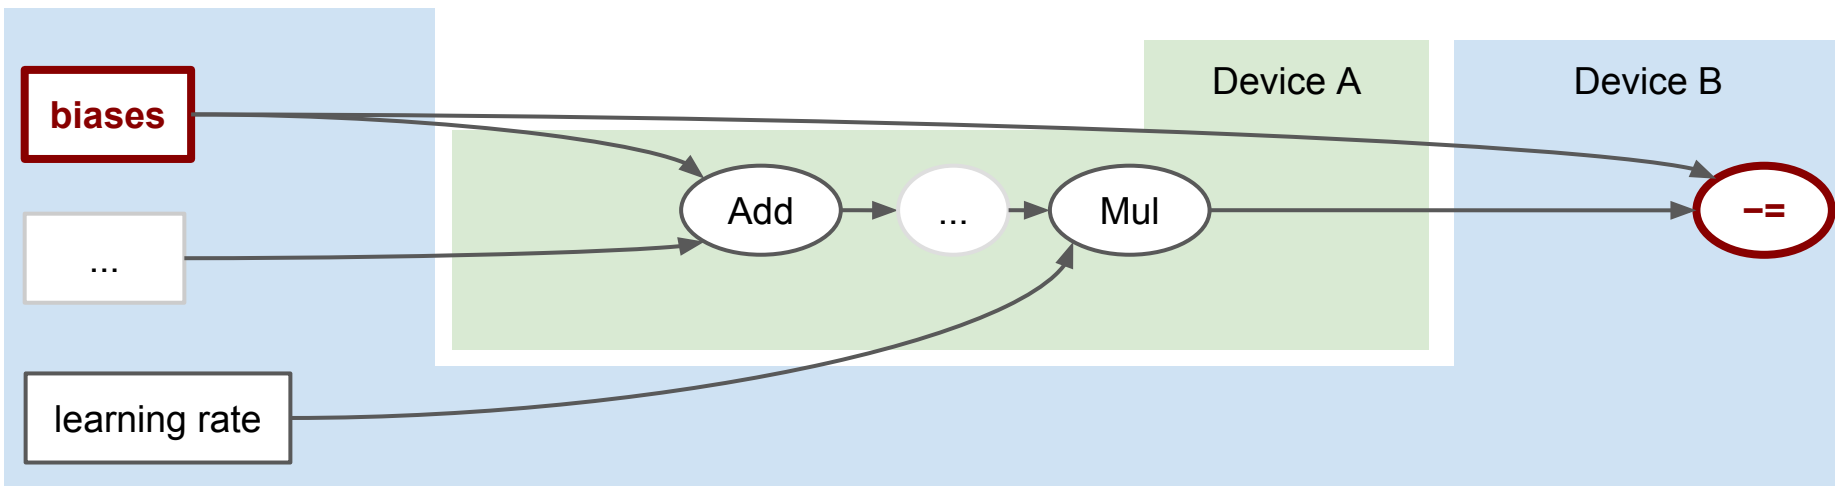
Simple gradient descent:

with state



Any Computation is a TensorFlow Graph


distributed



Devices: Processes, Machines, CPUs, GPUs, TPUs, etc

Send and Receive Nodes

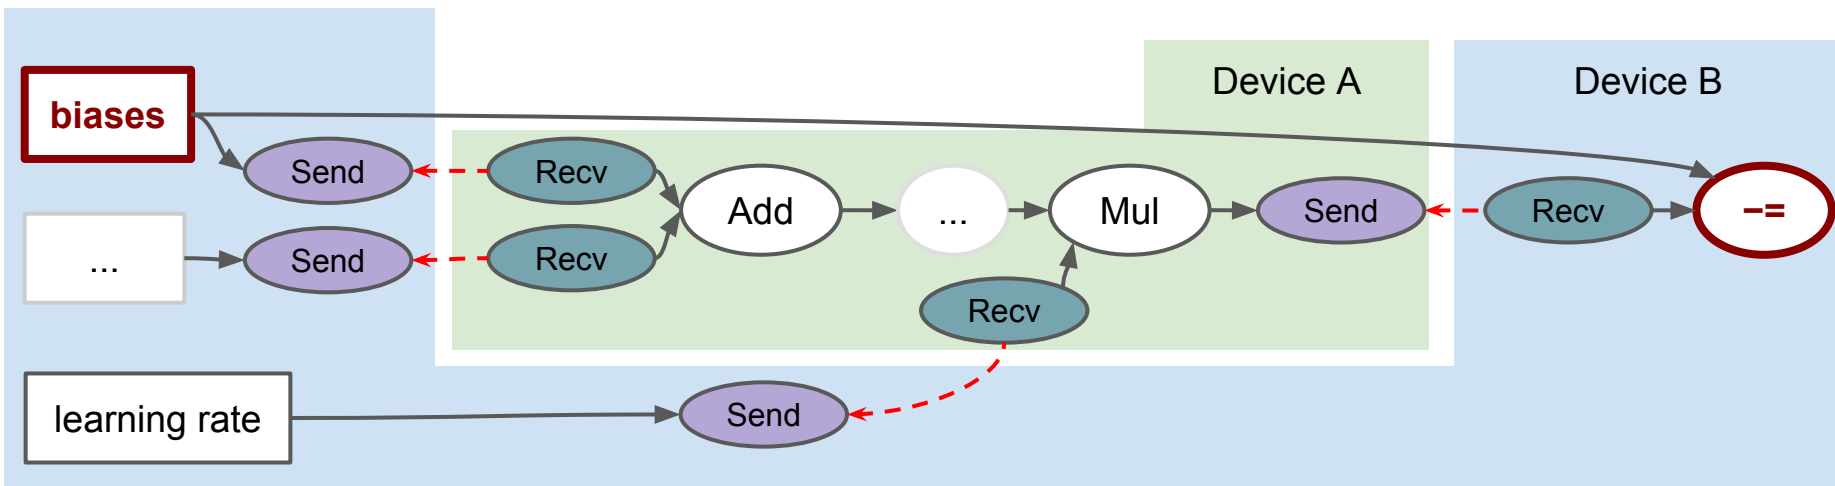
distributed



Devices: Processes, Machines, CPUs, GPUs, TPUs, etc

Send and Receive Nodes

distributed



Devices: Processes, Machines, CPUs, GPUs, TPUs, etc

Linear Regression


The background of the slide features a solid orange gradient. In the lower half, there is a stylized representation of a mountain range using a white wireframe mesh. The mesh lines are thin and white, creating a 3D effect against the orange background. The mountains are jagged and span the width of the slide.

Linear Regression

Diagram illustrating the linear regression equation $y = Wx + b$.

The equation components are labeled with arrows:

- result**: Points to y .
- input**: Points to x .
- parameters**: Points to W and b .



What are we trying to do?

Mystery equation: $y = 0.1 * x + 0.3 + \text{noise}$

Model: $y = W * x + b$

Objective: Given enough (x, y) value samples, figure out the value of W and b .

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf
```

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf
```

```
x = tf.placeholder(shape=[None],  
                   dtype=tf.float32, name="x")
```

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf
```

```
x = tf.placeholder(shape=[None],  
                    dtype=tf.float32, name="x")
```

```
W = tf.get_variable(shape=[], name="W")
```

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf
```

```
x = tf.placeholder(shape=[None],  
                    dtype=tf.float32, name="x")
```

```
W = tf.get_variable(shape=[], name="W")
```

```
b = tf.get_variable(shape=[], name="b")
```

$y = Wx + b$ in TensorFlow


```
import tensorflow as tf
```

```
x = tf.placeholder(shape=[None],  
                    dtype=tf.float32, name="x")
```

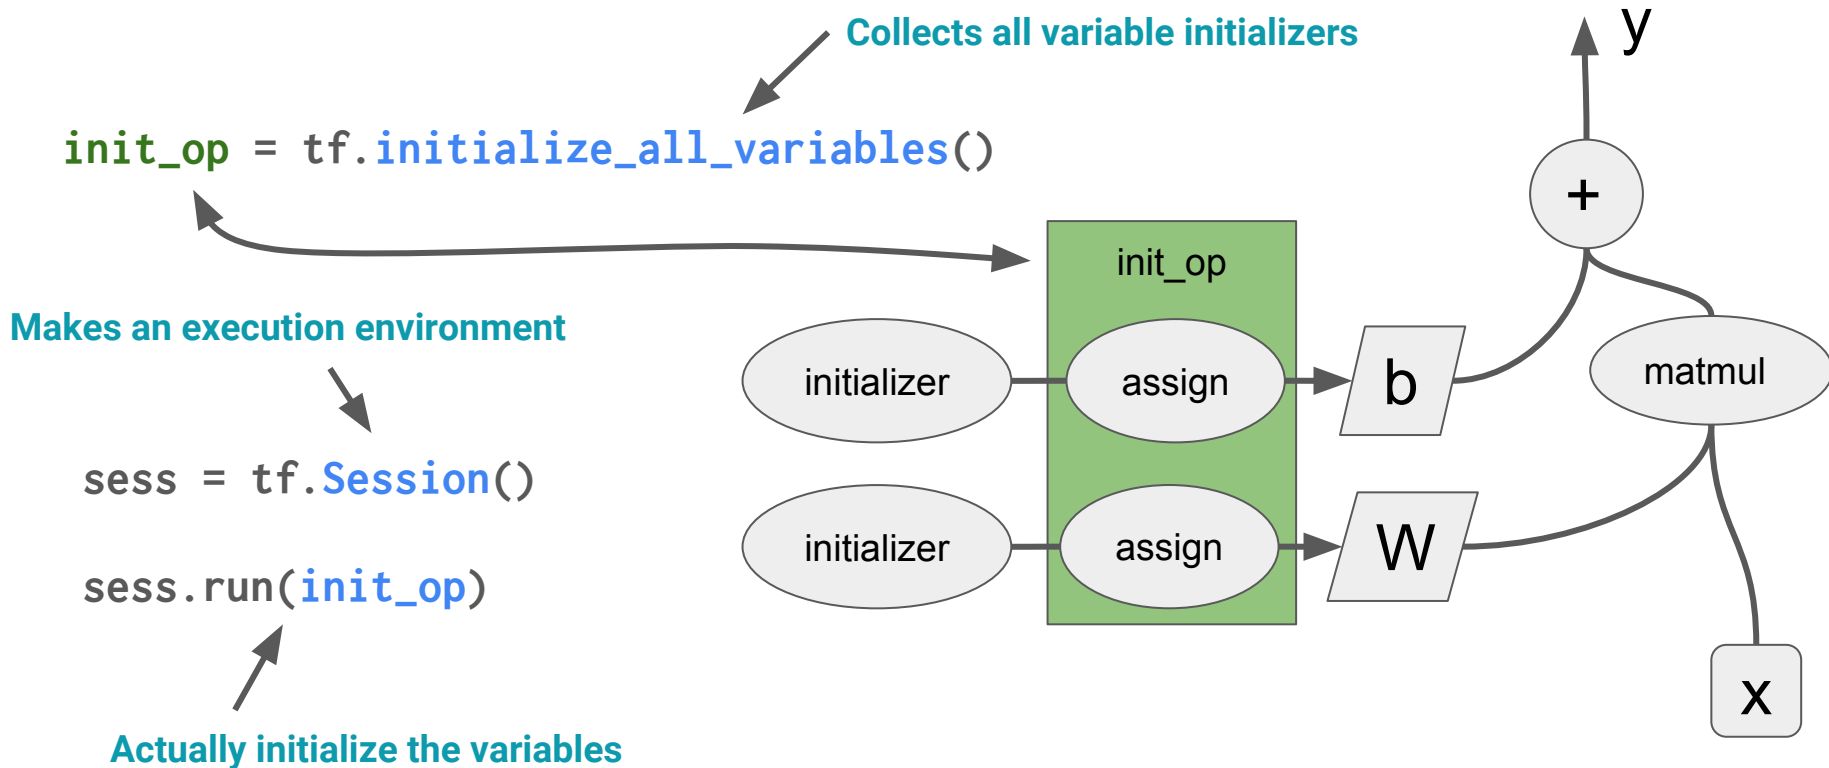
```
W = tf.get_variable(shape=[], name="W")
```

```
b = tf.get_variable(shape=[], name="b")
```

```
y = W * x + b
```



Variables Must be Initialized




Running the Computation

`x_in = 3`

`sess.run(y, feed_dict={x: x_in})`

- Only what's used to compute a fetch will be evaluated
- All Tensors can be fed, but all placeholders must be fed

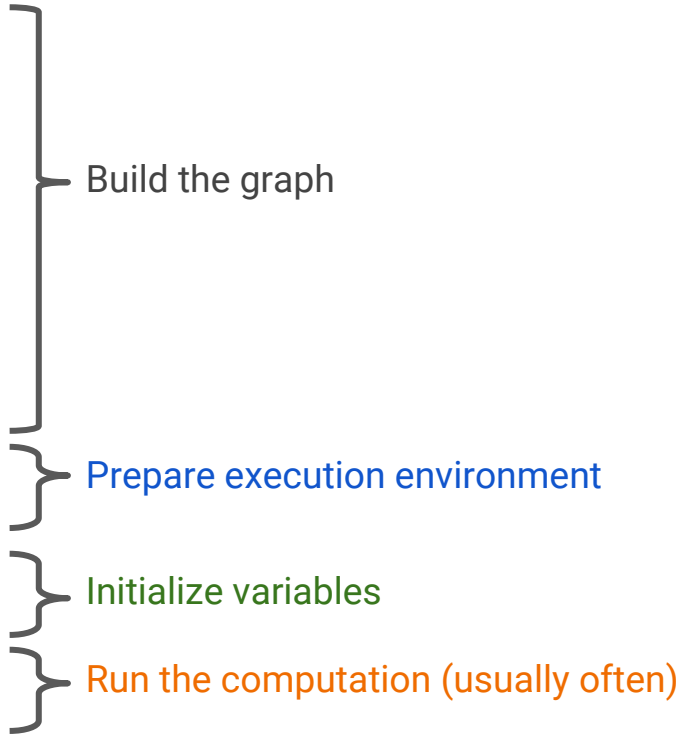


Putting it all together

```
import tensorflow as tf
x = tf.placeholder(shape=[None],
                   dtype=tf.float32,
                   name='x')

W = tf.get_variable(shape=[], name='W')
b = tf.get_variable(shape=[], name='b')
y = W * x + b

with tf.Session() as sess:
    sess.run(tf.initialize_all_variables())
    print(sess.run(y, feed_dict={x: x_in}))
```



The code is annotated with four steps, each indicated by a bracket on the right side:

- Build the graph**: This step covers the first five lines of code, which define the computational graph by creating a placeholder for input `x`, and variables `W` and `b`, and defining the operation `y = W * x + b`.
- Prepare execution environment**: This step covers the line `with tf.Session() as sess:`, which starts a new TensorFlow session.
- Initialize variables**: This step covers the line `sess.run(tf.initialize_all_variables())`, which initializes the variables `W` and `b`.
- Run the computation (usually often)**: This step covers the line `print(sess.run(y, feed_dict={x: x_in}))`, which runs the computation and prints the result.


Define a Loss

Given x , y compute a loss, for instance:

$$L = (y - y_{label})^2$$

```
# create an operation that calculates loss.  
loss = tf.reduce_mean(tf.square(y - y_data))
```

Minimize loss: optimizers



`tf.train.AdadeltaOptimizer`

`tf.train.AdagradOptimizer`

`tf.train.AdagradDAOptimizer`

`tf.train.AdamOptimizer`


...

Train

Feed (x, y_{label}) pairs and adjust W and b to decrease the loss.

$$W \leftarrow W - \eta (dL/dW)$$

$$b \leftarrow b - \eta (dL/db)$$



TensorFlow computes
gradients automatically

```
# Create an optimizer
```

```
optimizer = tf.train.GradientDescentOptimizer(0.5)
```

```
# Create an operation that minimizes loss.
```



Learning rate

```
train = optimizer.minimize(loss)
```

Putting it all together

```
loss = tf.reduce_mean(tf.square(y - y_label))
```

} Define a loss

```
optimizer = tf.train.GradientDescentOptimizer(0.5)
```

} Create an optimizer

```
train = optimizer.minimize(loss)
```

} Op to minimize the loss

```
with tf.Session() as sess:
```

```
    sess.run(tf.initialize_all_variables())
```

} Initialize variables

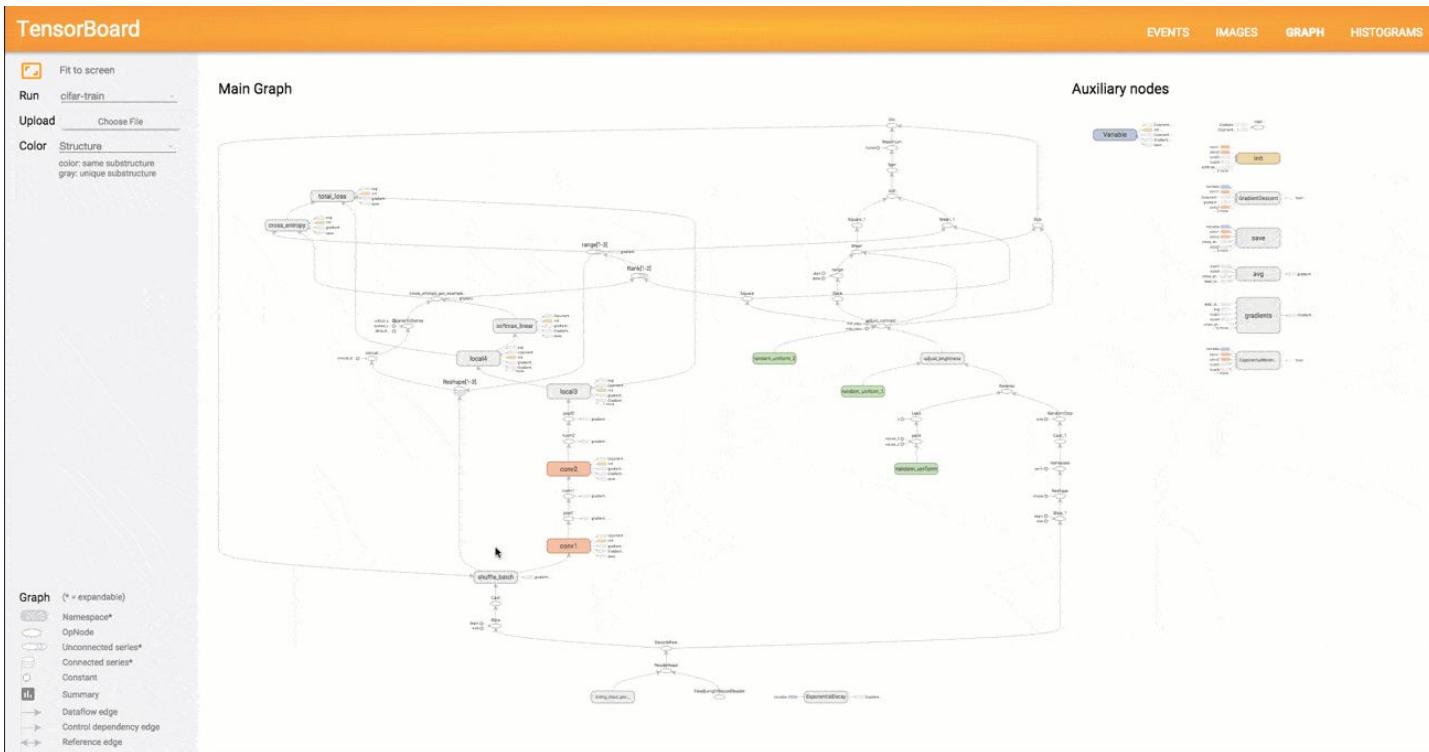
```
    for i in range(1000):
```

```
        sess.run(train, feed_dict={x: x_in[i],
```

```
                                   y_label: y_in[i]}))
```

} Iteratively run the training op

TensorBoard



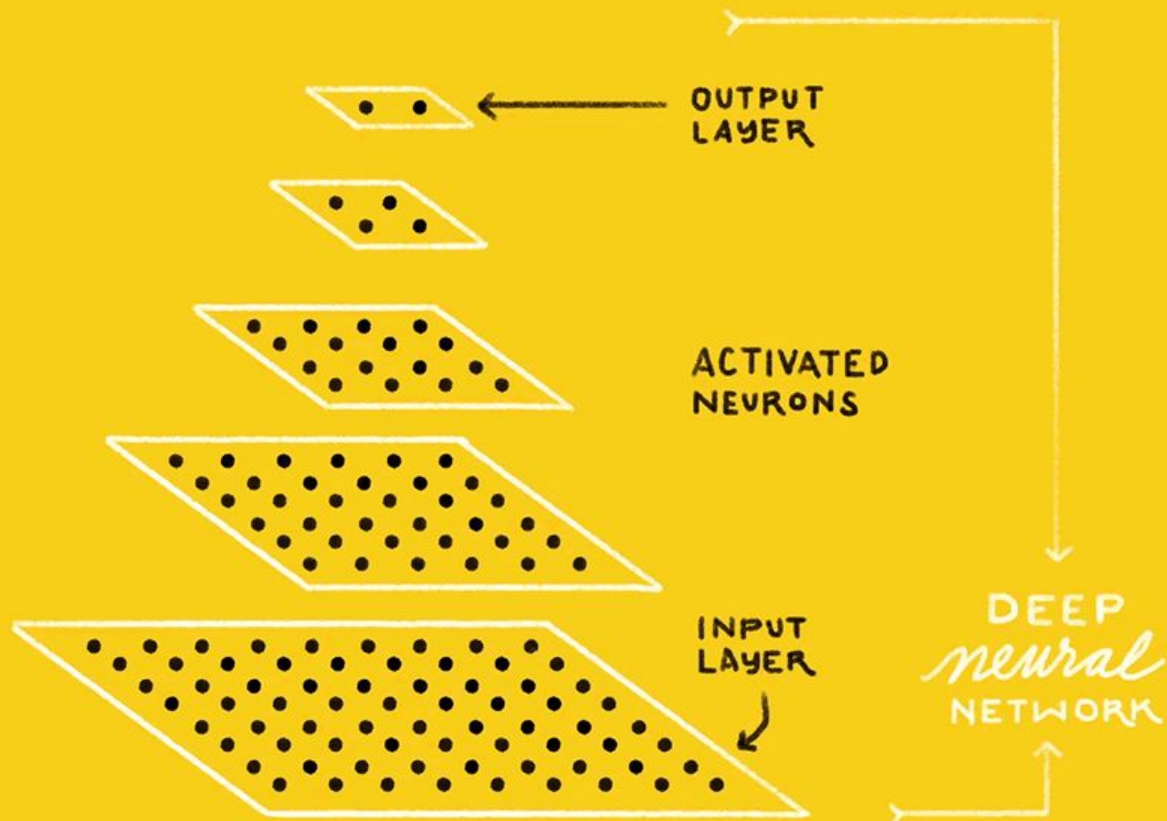
Deep Neural Network

The background of the slide features a solid orange gradient. Overlaid on this is a white, low-poly wireframe mesh that creates a sense of depth and texture, resembling a stylized landscape or a digital terrain. The mesh is composed of numerous interconnected lines forming a grid-like pattern that undulates across the frame.

IS THIS A
CAT or DOG?

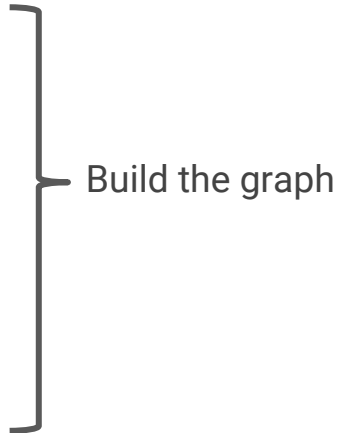


CAT DOG



Remember linear regression?


```
import tensorflow as tf
x = tf.placeholder(shape=[None],
                   dtype=tf.float32,
                   name='x')
W = tf.get_variable(shape=[], name='W')
b = tf.get_variable(shape=[], name='b')
y = W * x + b
loss = tf.reduce_mean(tf.square(y - y_label))
optimizer = tf.train.GradientDescentOptimizer(0.5)
train = optimizer.minimize(loss)
...
```




Build the graph

Convolutional DNN

```
x = tf.contrib.layers.conv2d(x, kernel_size=[5,5], ...)
x = tf.contrib.layers.max_pool2d(x, kernel_size=[2,2], ...)
x = tf.contrib.layers.conv2d(x, kernel_size=[5,5], ...)
x = tf.contrib.layers.max_pool2d(x, kernel_size=[2,2], ...)
x = tf.contrib.layers.fully_connected(x, activation_fn=tf.nn.relu)
x = tf.contrib.layers.dropout(x, 0.5)
logits = tf.config.layers.linear(x)
```




Defining Complex Networks



Distributed TensorFlow

The background of the slide features a solid orange gradient. Overlaid on this is a white wireframe pattern that forms a series of undulating, mountain-like shapes across the lower half of the image. The text 'Distributed TensorFlow' is centered horizontally and positioned in the upper-middle section of the slide.

Data Parallelism



Describe a cluster: ClusterSpec

```
tf.train.ClusterSpec({  
    "worker": [  
        "worker0.example.com:2222",  
        "worker1.example.com:2222",  
        "worker2.example.com:2222"  
    ],  
    "ps": [  
        "ps0.example.com:2222",  
        "ps1.example.com:2222"  
    ]  
})
```

Share the graph across devices

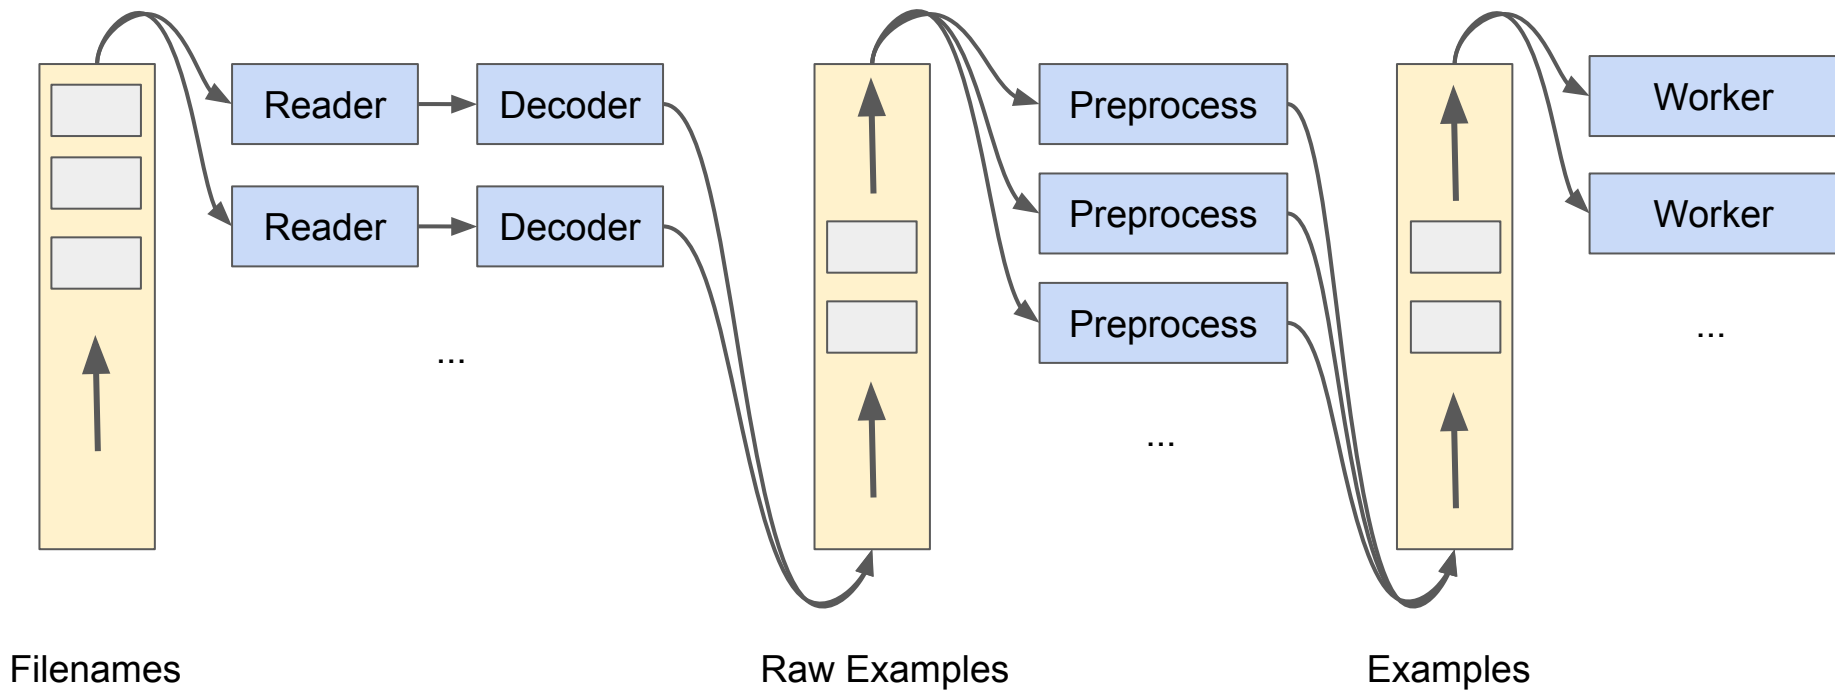
```
with tf.device("/job:ps/task:0"):
    weights_1 = tf.Variable(...)
    biases_1 = tf.Variable(...)
```

```
with tf.device("/job:ps/task:1"):
    weights_2 = tf.Variable(...)
    biases_2 = tf.Variable(...)
```

```
with tf.device("/job:worker/task:7"):
    input, labels = ...
    layer_1 = tf.nn.relu(tf.matmul(input, weights_1) + biases_1)
    logits = tf.nn.relu(tf.matmul(layer_1, weights_2) + biases_2)
    train_op = ...
```

```
with tf.Session("grpc://worker7.example.com:2222") as sess:
    for _ in range(10000):
        sess.run(train_op)
```


Input Pipelines with Queues



Tutorials & Courses

Tutorials on [tensorflow.org](https://www.tensorflow.org):

Image recognition: https://www.tensorflow.org/tutorials/image_recognition

Word embeddings: <https://www.tensorflow.org/versions/word2vec>

Language Modeling: <https://www.tensorflow.org/tutorials/recurrent>

Translation: <https://www.tensorflow.org/versions/seq2seq>

Deep Dream:

<https://tensorflow.org/code/tensorflow/examples/tutorials/deepdream/deepdream.ipynb>

Thank you and have fun!



Martin Wicke
[@martin_wicke](#)



Rajat Monga
[@rajatmonga](#)

The background is a solid orange color. In the lower half, there is a white wireframe illustration of a mountain range. The mountains are composed of a grid of lines that form a series of peaks and valleys, creating a three-dimensional effect. The word "Extras" is written in a bold, white, sans-serif font, positioned on the left side of the image, overlapping the wireframe mountains.

Extras


Inception



An Alaskan Malamute (left) and a Siberian Husky (right). Images from Wikipedia.

<https://research.googleblog.com/2016/08/improving-inception-and-image.html>


Show and Tell



A person on a beach flying a kite.

<https://research.googleblog.com/2016/09/show-and-tell-image-captioning-open.html>

Parsey McParseface



<https://research.googleblog.com/2016/05/announcing-syntaxnet-worlds-most.html>

Text Summarization

Original text

- *Alice and Bob took the train to visit the zoo. They saw a **baby giraffe, a lion, and a flock of colorful tropical birds.***

Abstractive summary

- *Alice and Bob visited the zoo and saw **animals and birds.***

Mobile TensorFlow

TensorFlow was designed with mobile and embedded platforms in mind. We have sample code and build support you can try now for these platforms:

[Android](#)

[iOS](#)

[Raspberry Pi](#)

Many applications can benefit from on-device processing. Google Translate's instant visual translation is a great example. By running its processing locally, users get an incredibly responsive and interactive experience.




Mobile TensorFlow makes sense when there is a poor or missing network connection, or where sending continuous data to a server would be too expensive. We are working to help developers make lean mobile apps using TensorFlow, both by continuing to reduce the code footprint, and supporting [quantization](#) and [lower precision arithmetic](#) that reduce model size.





Monet - Sunflowers: 0.992





Architecture

